

Instruction Set Architecture

(Instruction encoding)

Acknowledgment: Slides are adapted from Harris and Harris
textbook instructor's material

Machine Language and Instruction Encoding

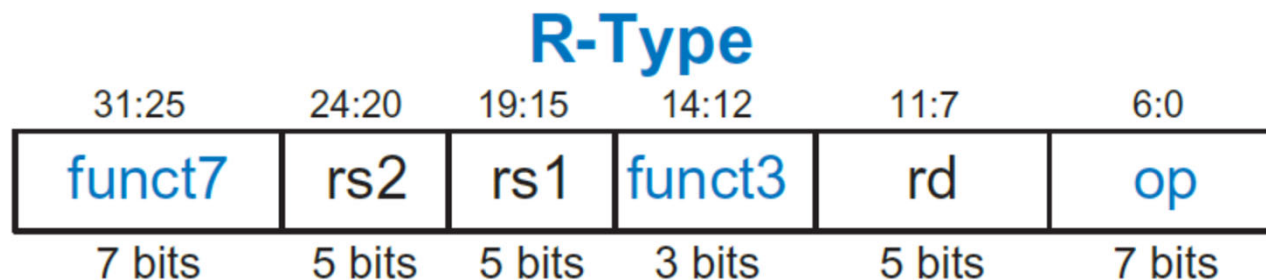
- Binary representation of instructions (computers only understand 1's and 0's)
- 32-bit instructions (for studied, baseline RV32)
 - Instruction bits are divided into fields, each field tells processor something about instruction
- Simplicity favors regularity → Six basic instruction formats (as opposed to, e.g., unique format for each instruction)
 - **R-Type**: for register-register arithmetic operations
 - **I-Type**: for register-immediate arithmetic operations and loads
 - **S-Type**: for stores
 - **B-Type**: for branches (similar to S-type)
 - **U-Type**: for 20-bit upper immediate instructions
 - **J-Type**: for jumping (similar to U-type)

RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0		
funct7				rs2			rs1		funct3		rd			opcode		R-type
imm[11:0]							rs1		funct3		rd			opcode		I-type
imm[11:5]				rs2			rs1		funct3		imm[4:0]			opcode		S-type
imm[12]	imm[10:5]			rs2			rs1		funct3		imm[4:1]	imm[11]		opcode	B-type	
imm[31:12]										rd			opcode		U-type	
imm[20]	imm[10:1]				imm[11]		imm[19:12]			rd			opcode		J-type	

R-Type (Register) Instruction Layout

- 3 register operands:
 - `rs1, rs2`: specify registers containing first and second operand
 - `rd`: specify register which will receive the results of computation
 - Each register field is 5-bit unsigned number encoding register number
- Other fields:
 - `opcode`: partially specifies what instruction does
 - `funct7+funct3`: combined with `opcode`, these fields describe the operation



R-Type Examples

Assembly

```
add s2, s3, s4
add x18, x19, x20
sub t0, t1, t2
sub x5, x6, x7
```

Field Values

funct7	rs2	rs1	funct3	rd	op
0	20	19	0	18	51
32	7	6	0	5	51
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits


Machine Code

funct7	rs2	rs1	funct3	rd	op	
0000,000	10100	10011	000	10010	011,0011	(0x01498933)
0100,000	00111	00110	000	00101	011,0011	(0x407302B3)
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	

Note the different order of registers in the assembly code

I-Type (Immediate) Instruction Layout

- Used for arithmetic / logical operations with immediate and load instructions
- Instruction fields:
 - `rs1, rd`: source and destination register operands (note that only two registers!)
 - `imm`: 12-bit two's complement immediate, i.e. any number from $[-2048_{\text{ten}}, +2047_{\text{ten}}]$, always sign-extended to 32-bits before using in arithmetic operation
 - `op`: the opcode completely determining operation for arithmetic operations; together with `funct3` field determines operation for others
- Load instructions are also I-Type
 - 12-bit immediate is used to create address, not the final result
- The format is similar R-Type to simplify instruction decoding



I-Type Examples

Assembly

Field Values

Machine Code

```
addi s0, s1, 12
addi x8, x9, 12
addi s2, t1, -14
addi x18, x6, -14
lw t2, -6(s3)
lw x7, -6(x19)
lb s4, 0x1F(s4)
lb x20, 0x1F(x20)
slli s2, s7, 5
slli x18, x23, 5
srai t1, t2, 29
srai x6, x7, 29
```

imm _{11:0}	rs1	funct3	rd	op
12	9	0	8	19
-14	6	0	18	19
-6	19	2	7	3
0x1F	20	0	20	3
5	23	1	18	19
(upper 7 bits = 32) 29	7	5	6	19
12 bits	5 bits	3 bits	5 bits	7 bits

imm _{11:0}	rs1	funct3	rd	op	
0000 0000 1100	01001	000	01000	001 0011	(0x00C48413)
1111 1111 0010	00110	000	10010	001 0011	(0xFF230913)
1111 1111 1010	10011	010	00111	000 0011	(0xFFA9A383)
0000 0001 1111	10100	000	10100	000 0011	(0x01FA0A03)
0000 0000 0101	10111	001	10010	001 0011	(0x005B9913)
0100 0001 1101	00111	101	00110	001 0011	(0x41D3D313)
12 bits	5 bits	3 bits	5 bits	7 bits	

Note again different order of registers in assembly and machine codes!

All RV32 I-type Arithmetic/Logical/Load Instructions

imm[11:0]		rs1	000	rd	0010011	ADDI
imm[11:0]		rs1	010	rd	0010011	SLTI
imm[11:0]		rs1	011	rd	0010011	SLTIU
imm[11:0]		rs1	100	rd	0010011	XORI
imm[11:0]		rs1	110	rd	0010011	ORI
imm[11:0]		rs1	111	rd	0010011	ANDI
0000000	shamt	rs1	001	rd	0010011	SLLI
0000000	shamt	rs1	101	rd	0010011	SRLI
0000000	shamt	rs1	101	rd	0010011	SRAI

One of the higher-order immediate bits is used to distinguish “shift right logical” (SRLI) from “shift right arithmetic” (SRAI)

“Shift-by-immediate” instructions only use lower 5 bits of the immediate value for shift amount (can only shift by 0-31 bit positions)

imm[11:0]	rs1	000	rd	0000011	LB
imm[11:0]	rs1	001	rd	0000011	LH
imm[11:0]	rs1	010	rd	0000011	LW
imm[11:0]	rs1	100	rd	0000011	LBU
imm[11:0]	rs1	101	rd	0000011	LHU

funct3 field encodes size and signedness of load data

S/B-Type Instruction Layout

- Used for stores and branches
- Instruction fields:
 - `rs1, rs2`: source1 and source2 register operands (none of them is written)
 - `imm`: 12-bit two's complement immediate, always sign-extended to 32-bits before using in arithmetic operation (i.e., address calculation or comparison)
 - `Op` and `funct3`: completely determines operation
- For performance reasons, `rs1` and `rs2` field locations (but not `imm`) are consistent with R-format

31:25	24:20	19:15	14:12	11:7	6:0	
imm _{11:5}	rs2	rs1	funct3	imm _{4:0}	op	S-Type
imm _{12,10:5}	rs2	rs1	funct3	imm _{4:1,11}	op	B-Type
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	

S-Type Example

Assembly

```
sw t2, -6(s3)
sw x7, -6(x19)
sh s4, 23(t0)
sh x20, 23(x5)
sb t5, 0x2D(zero)
sb x30, 0x2D(x0)
```

Field Values

imm _{11:5}	rs2	rs1	funct3	imm _{4:0}	op
1111 111	7	19	2	11010	35
0000 000	20	5	1	10111	35
0000 001	30	0	0	01101	35
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

Machine Code

imm _{11:5}	rs2	rs1	funct3	imm _{4:0}	op	
1111 111	00111	10011	010	11010	010 0011	(0xFE79AD23)
0000 000	10100	00101	001	10111	010 0011	(0x01429BA3)
0000 001	11110	00000	000	01101	010 0011	(0x03E006A3)
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	

PC-Relative Addressing

- Use the immediate field as a two's complement offset to PC
 - Branches generally change the PC by small amount
 - To improve the reach of a single branch instruction, we can encode word offset rather than a byte offset (typical for other RISC architectures)
 - However for compatibility reasons with 16-bit compressed instructions extension, encode half word offsets
 - Can jump to $\pm 2^{10}$ 32-bit instructions from the current PC
- Branch calculation scheme:
 - Need to multiply the offset by 2 before adding to PC
 - $PC = PC + \text{immediate} * 2$
 - Immediate is a number of instructions to jump either forward or backwards

Branch Example

#Address	#	RISC-V	Assembly
0x70		beq	s0, t5, L1
0x74		add	s1, s2, s3
0x78		sub	s5, s6, s7
0x7C		lw	t0, 0(s1)
0x80	L1:	addi	s1, s1, -15


L1 is 4 instructions (i.e., **16 bytes**) past beq

imm _{12:0} = 16	0	0	0	0	0	0	0	0	1	0	0	0	X
bit number	12	11	10	9	8	7	6	5	4	3	2	1	0

Assembly	Field Values						Machine Code						
beq s0, t5, L1 beq x8, x30, 16	imm _{12,10:5}	rs2	rs1	funct3	imm _{4:1,11}	op	imm _{12,10:5}	rs2	rs1	funct3	imm _{4:1,11}	op	
	0000 000	30	8	0	1000 0	99	0000 000	11110	01000	000	1000 0	110 0011	(0x01E40863)
	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	

U/J-Type Instruction Layout

- Used for loading 20-bit immediate and unconditional branches
- Only one destination register
- opcode completely determines operation



- J-Type: JAL instruction
 - JAL saves PC+4 in register rd (the return address)
 - Set PC to PC + immediate * 2
 - PC-relative addressing with $\pm 2^{18}$ 32-bit instruction range from the current value of PC (using similar trick to B-type format)
- Assembler “J” jump is pseudo-instruction, uses JAL with rd=x0 to discard return address

Jump Example

```
# Address      RISC-V Assembly
0x0000540C    jal ra, func1
0x00005410    add s1, s2, s3
...           ...

0x000ABC04    func1: add s4, s5, s8
...           ...

func1 is 0xA67F8 bytes past jal
```


imm = 0xA67F8 0 1 0 1 0 0 1 1 0 0 1 1 1 1 1 1 1 0 0 0

bit number 20 19 18 17 16 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0

Assembly	Field Values			Machine Code			
	imm _{20,10:1,11,19:12}	rd	op	imm _{20,10:1,11,19:12}	rd	op	
jal ra, func1 jal x1, 0xA67F8	0111 1111 1000 1010 0110	1	111	0111 1111 1000 1010 0110	00001	110 1111	(0x7F8A60EF)
	20 bits	5 bits	7 bits	20 bits	5 bits	7 bits	

U-Type Instructions

- U-type instructions
 - LUI – load upper immediate (writes the upper 20 bits of the destination with the immediate values, and clears the lower 12 bits)
 - AUIPC – to be discussed shortly



U-type example:

Assembly

Field Values

Machine Code

	imm _{31:12}	rd	op				
lui s5, 0x8CDEF	0x8CDEF	21	55	1000 1100 1101 1110 1111	10101	011 0111	(0x8CDEFAB7)
lui x21, 0x8CDEF							
	20 bits	5 bits	7 bits	20 bits	5 bits	7 bits	

Loading Large Constants Revisited (Corner Case)

- Recollect how 32 bit constants are loaded (e.g., 0xDEADBEEF → x10)
 - LUI x10, 0xDEADB
 - ADDI x10, x10, 0xEEF
- However, ADDI 12-bit immediate is always sign-extended, so if the top bit is 1, it will add “-1” to upper 20 bits
 - Instead of 0xDEADBEEF, it will be 0xDEADAEEF
- Solution: Pre-increment value placed in upper 20 bits, if the top bit is set in the immediate lower 12 bits
- Assembler pseudo-instruction LI takes care of it
 - LI x10, 0xDEADBEEF → LUI x10, 0xDEADC followed by ADDI x10,x10, 0xEEF

Far-Away Jumps

- JAL (J) and branches jump relative to PC-relative. The range is limited (up to $\pm 2^{18}$ instructions for JALs)
- For far away branches (i.e. call to O/S functions), use combination of AUIPC and JALR instructions
- AUIPC (U-type instruction): Add 20-bit immediate to upper PC bits and places the result in the destination register
 - AUIPC rd, upimm $rd = \{upimm, 12'b0\} + PC$
 - Label: AUIPC x10, 0 # puts address of label in x10
- JALR (I-type instruction): JALR rd, rs, immediate
 - writes PC+4 to rd
 - Sets $PC = rs + immediate$ (note: no multiplication by 2)

Far-away jump example:

Address of instruction specified by label L5

call L5

auipc ra, imm_{31:12}

jalr ra, ra, imm_{11:0}

call far away function $PC = L5, ra = PC + 4$

Immediate Encoding Summary

Instruction encodings (31:7 bits only):

11	10	9	8	7	6	5	4	3	2	1	0	rs1	funct3	rd										
11	10	9	8	7	6	5	rs2					rs1	funct3	4	3	2	1	0						
12	10	9	8	7	6	5	rs2					rs1	funct3	4	3	2	1	11						
31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	rd				
20	10	9	8	7	6	5	4	3	2	1	11	19	18	17	16	15	14	13	12	rd				
31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7

32-bit immediate produced:

imm ₁₁												imm _{11:1}								imm ₀		I, B U J									
imm ₁₂												imm _{11:1}								0											
imm _{31:21}						imm _{20:12}						0																			
imm ₂₀						imm _{20:12}						imm _{11:1}								0											
31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0

RISC-V Base Instruction Reference Sheet

- 2nd page of Harris & Harris book

31:25		24:20		19:15	14:12	11:7	6:0	
funct7		rs2	rs1	funct3	rd	op		R-Type
imm _{11:0}			rs1	funct3	rd	op		I-Type
imm _{11:5}		rs2	rs1	funct3	imm _{4:0}	op		S-Type
imm _{12,10:5}		rs2	rs1	funct3	imm _{4:1,11}	op		B-Type
imm _{31:12}					rd	op		U-Type
imm _{20,10:1,11,19:12}					rd	op		J-Type
fs3	funct2	fs2	fs1	funct3	fd	op		R4-Type
5 bits		2 bits	5 bits	5 bits	3 bits	5 bits	7 bits	

- imm: signed immediate in imm_{11:0}
- uimm: 5-bit unsigned immediate in imm_{4:0}
- upimm: 20 upper bits of a 32-bit immediate, in imm_{31:12}
- Address: memory address: rs1 + SignExt(imm_{11:0})
- [Address]: data at memory location Address
- BTA: branch target address: PC + SignExt((imm_{12:1}, 1'b0))
- JTA: jump target address: PC + SignExt((imm_{20:1}, 1'b0))
- label: text indicating instruction address
- SignExt: value sign-extended to 32 bits
- ZeroExt: value zero-extended to 32 bits
- csr: control and status register

op	funct3	funct7	Type	Instruction	Description	Operation
0000011 (3)	000	-	I	lb rd, imm(rs1)	load byte	rd = SignExt([Address] _{7:0})
0000011 (3)	001	-	I	lh rd, imm(rs1)	load half	rd = SignExt([Address] _{15:0})
0000011 (3)	010	-	I	lw rd, imm(rs1)	load word	rd = [Address] _{31:0}
0000011 (3)	100	-	I	lbu rd, imm(rs1)	load byte unsigned	rd = ZeroExt([Address] _{7:0})
0000011 (3)	101	-	I	lhu rd, imm(rs1)	load half unsigned	rd = ZeroExt([Address] _{15:0})
0010011 (19)	000	-	I	addi rd, rs1, imm	add immediate	rd = rs1 + SignExt(imm)
0010011 (19)	001	0000000*	I	slli rd, rs1, uimm	shift left logical immediate	rd = rs1 << uimm
0010011 (19)	010	-	I	slti rd, rs1, imm	set less than immediate	rd = (rs1 < SignExt(imm))
0010011 (19)	011	-	I	sltiu rd, rs1, imm	set less than imm. unsigned	rd = (rs1 < SignExt(imm))
0010011 (19)	100	-	I	xori rd, rs1, imm	xor immediate	rd = rs1 ^ SignExt(imm)
0010011 (19)	101	0000000*	I	srli rd, rs1, uimm	shift right logical immediate	rd = rs1 >> uimm
0010011 (19)	101	0100000*	I	srai rd, rs1, uimm	shift right arithmetic imm.	rd = rs1 >>> uimm
0010011 (19)	110	-	I	ori rd, rs1, imm	or immediate	rd = rs1 SignExt(imm)
0010011 (19)	111	-	I	andi rd, rs1, imm	and immediate	rd = rs1 & SignExt(imm)
0010111 (23)	-	-	U	auipc rd, upimm	add upper immediate to PC	rd = {upimm, 12'b0} + PC
0100011 (35)	000	-	S	sb rs2, imm(rs1)	store byte	[Address] _{7:0} = rs2 _{7:0}
0100011 (35)	001	-	S	sh rs2, imm(rs1)	store half	[Address] _{15:0} = rs2 _{15:0}
0100011 (35)	010	-	S	sw rs2, imm(rs1)	store word	[Address] _{31:0} = rs2
0110011 (51)	000	0000000	R	add rd, rs1, rs2	add	rd = rs1 + rs2
0110011 (51)	000	0100000	R	sub rd, rs1, rs2	sub	rd = rs1 - rs2
0110011 (51)	001	0000000	R	sll rd, rs1, rs2	shift left logical	rd = rs1 << rs2 _{4:0}
0110011 (51)	010	0000000	R	slt rd, rs1, rs2	set less than	rd = (rs1 < rs2)
0110011 (51)	011	0000000	R	sltu rd, rs1, rs2	set less than unsigned	rd = (rs1 < rs2)
0110011 (51)	100	0000000	R	xor rd, rs1, rs2	xor	rd = rs1 ^ rs2
0110011 (51)	101	0000000	R	srl rd, rs1, rs2	shift right logical	rd = rs1 >> rs2 _{4:0}
0110011 (51)	101	0100000	R	sra rd, rs1, rs2	shift right arithmetic	rd = rs1 >>> rs2 _{4:0}
0110011 (51)	110	0000000	R	or rd, rs1, rs2	or	rd = rs1 rs2
0110011 (51)	111	0000000	R	and rd, rs1, rs2	and	rd = rs1 & rs2
0110111 (55)	-	-	U	lui rd, upimm	load upper immediate	rd = {upimm, 12'b0}
1100011 (99)	000	-	B	beq rs1, rs2, label	branch if =	if (rs1 == rs2) PC = BTA
1100011 (99)	001	-	B	bne rs1, rs2, label	branch if ≠	if (rs1 ≠ rs2) PC = BTA
1100011 (99)	100	-	B	blt rs1, rs2, label	branch if <	if (rs1 < rs2) PC = BTA
1100011 (99)	101	-	B	bge rs1, rs2, label	branch if ≥	if (rs1 ≥ rs2) PC = BTA
1100011 (99)	110	-	B	bltu rs1, rs2, label	branch if < unsigned	if (rs1 < rs2) PC = BTA
1100011 (99)	111	-	B	bgeu rs1, rs2, label	branch if ≥ unsigned	if (rs1 ≥ rs2) PC = BTA
1100111 (103)	000	-	I	jalr rd, rs1, imm	jump and link register	PC = rs1 + SignExt(imm), rd = PC + 4
1101111 (111)	-	-	J	jal rd, label	jump and link	PC = JTA, rd = PC + 4

*Encoded in instr_{31:25}, the upper seven bits of the immediate field